Requesting GPU Resources on Google Cloud Platform

1. Navigate to the **Quotas** page under **IAM & Admin** section.

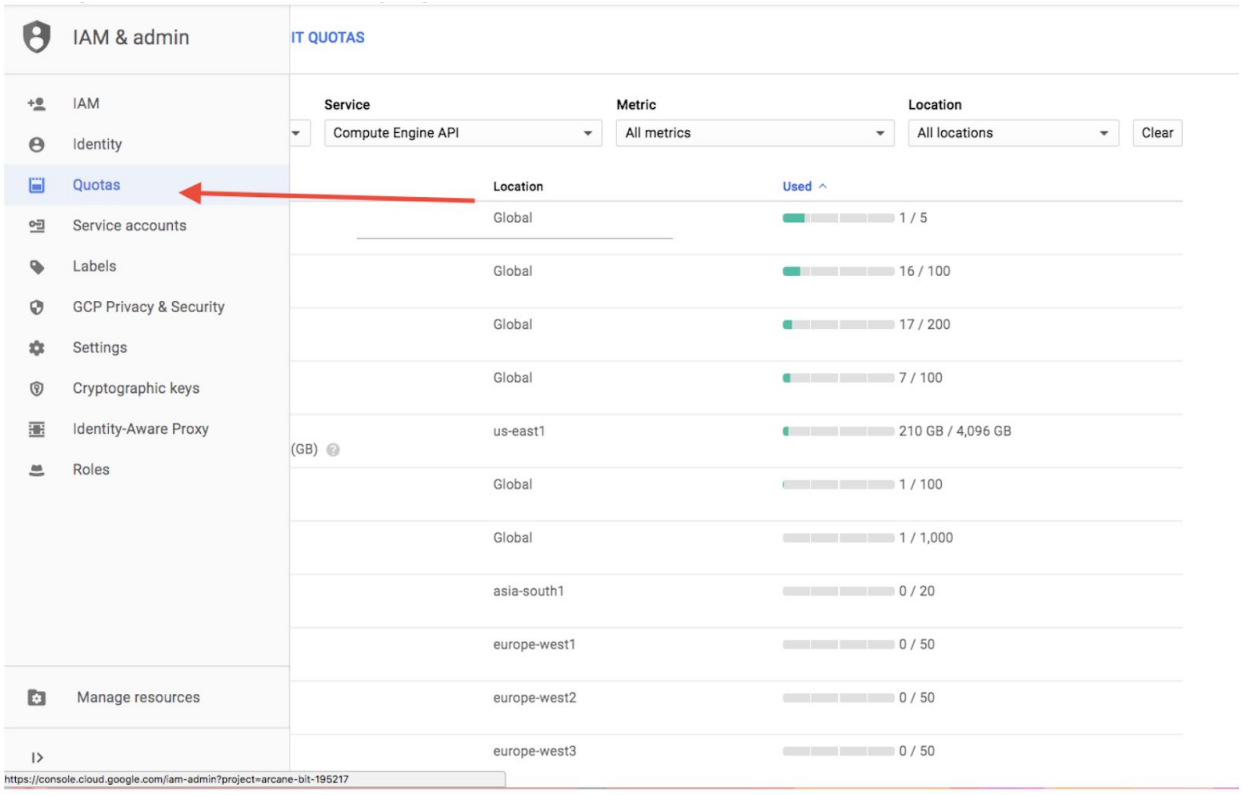

2. Under the **Metric** window, select None.

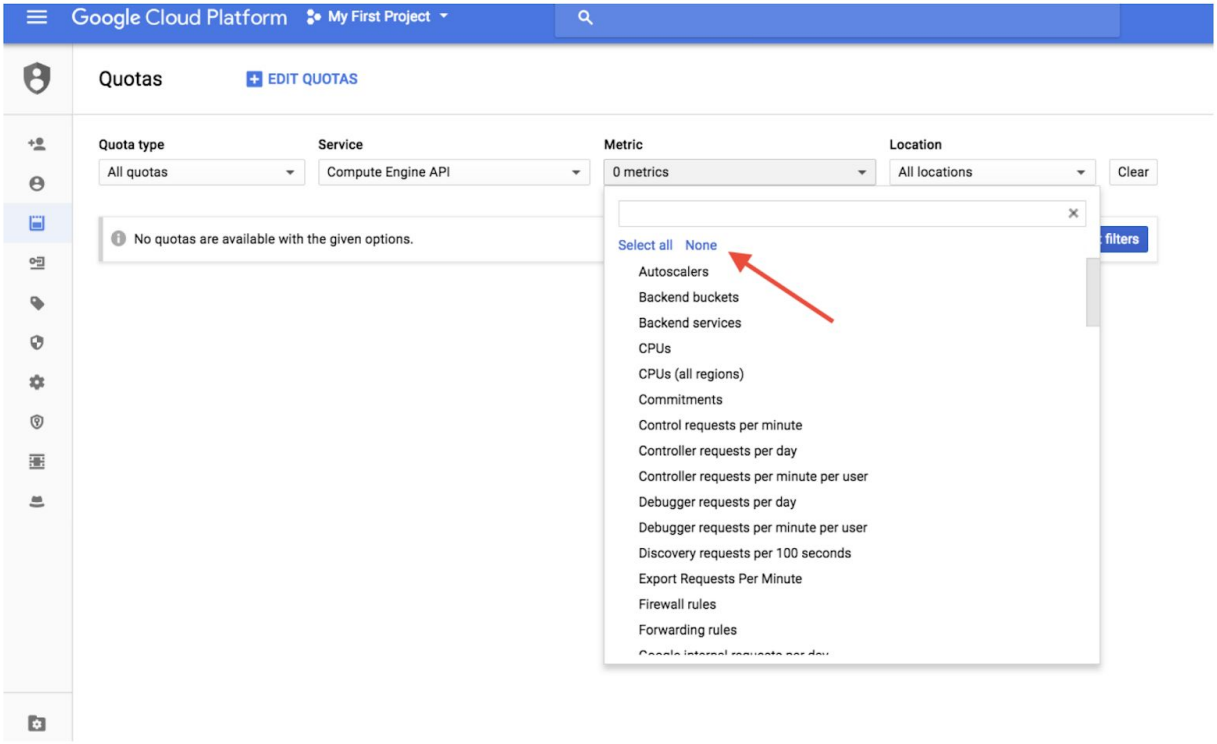

3. Type ‘gpu’ into the search box and select **NVIDIA K80 GPUs** with location **us-east-1**.

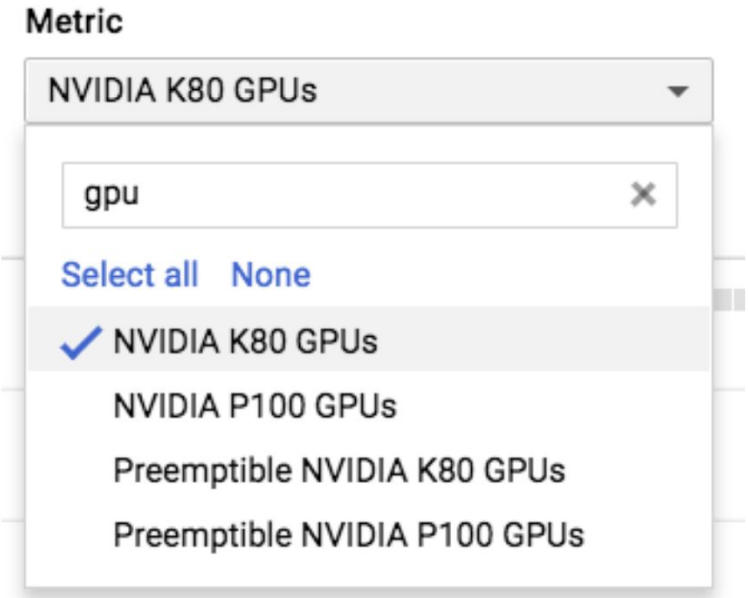

4. Select **Edit Quotas**

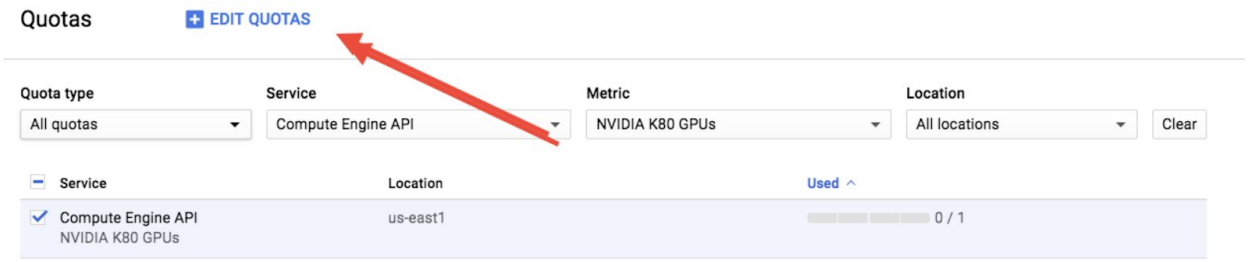

5. Fill in your name, email and phone number, then fill in the following form:

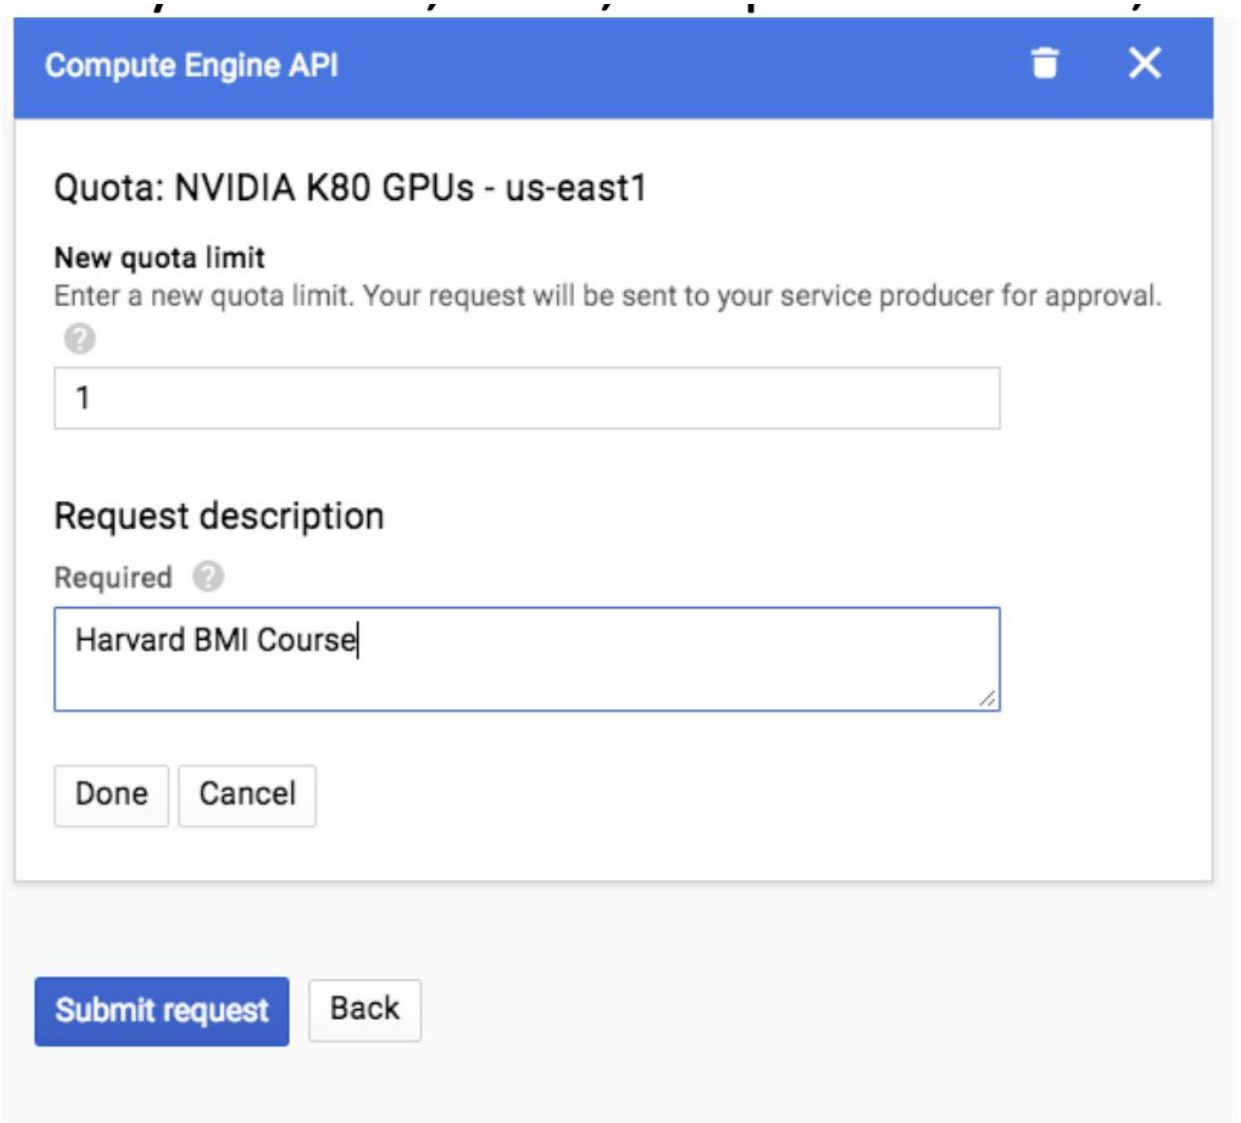

6. Wait ~5-10 minutes for approval.
7. When creating a VM, ensure the Zone is set to **us-east1-c** and that a **GPU** is selected. Only certain zones can run a GPU.

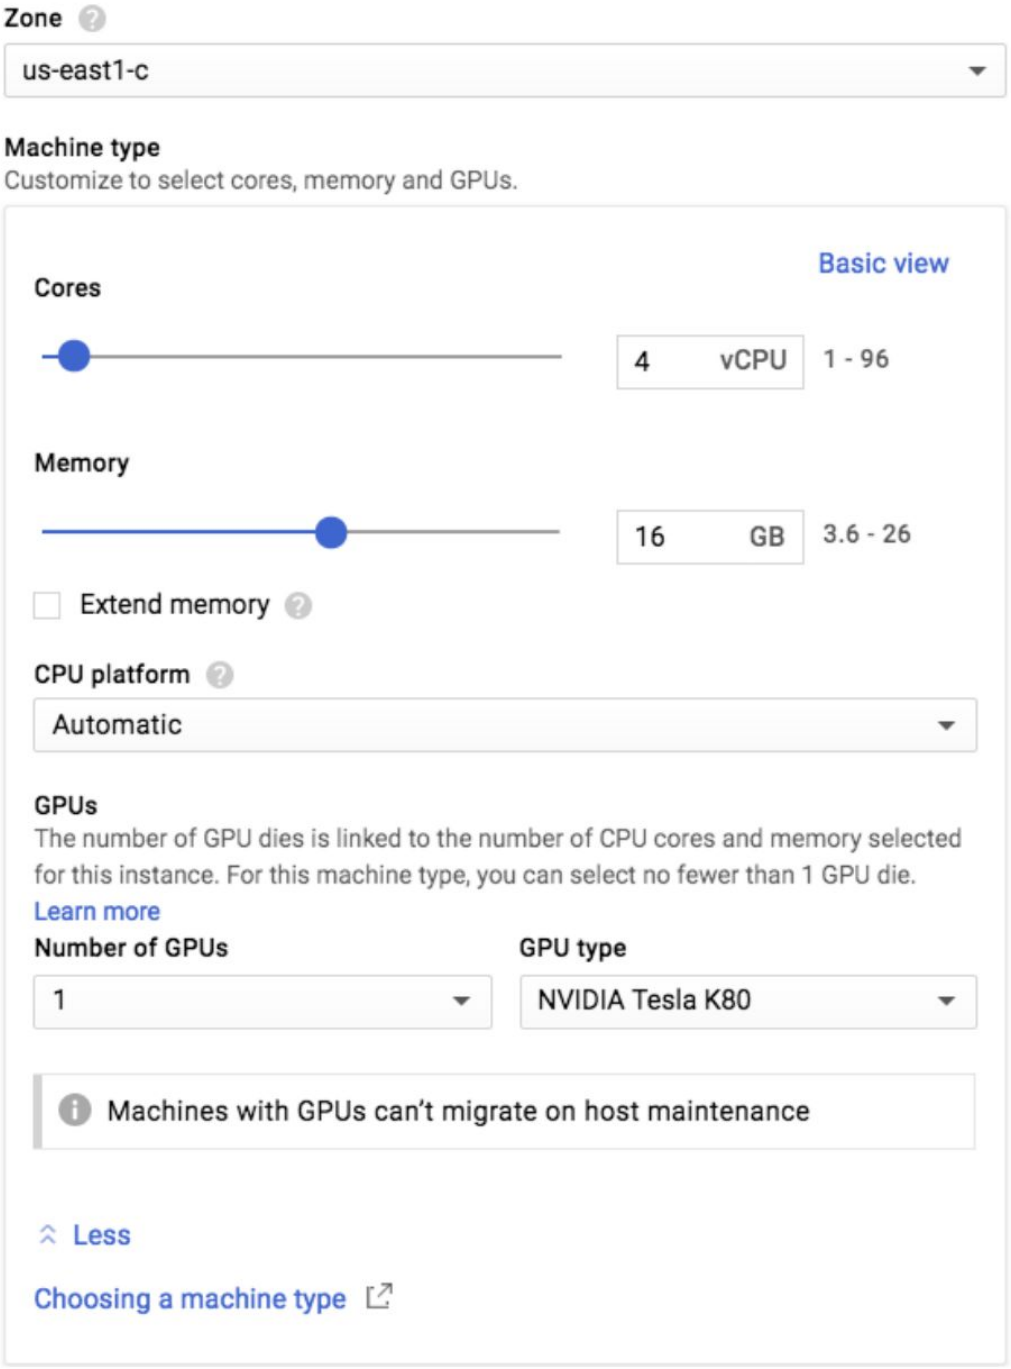

8. Under **Boot Options**, select the following options:

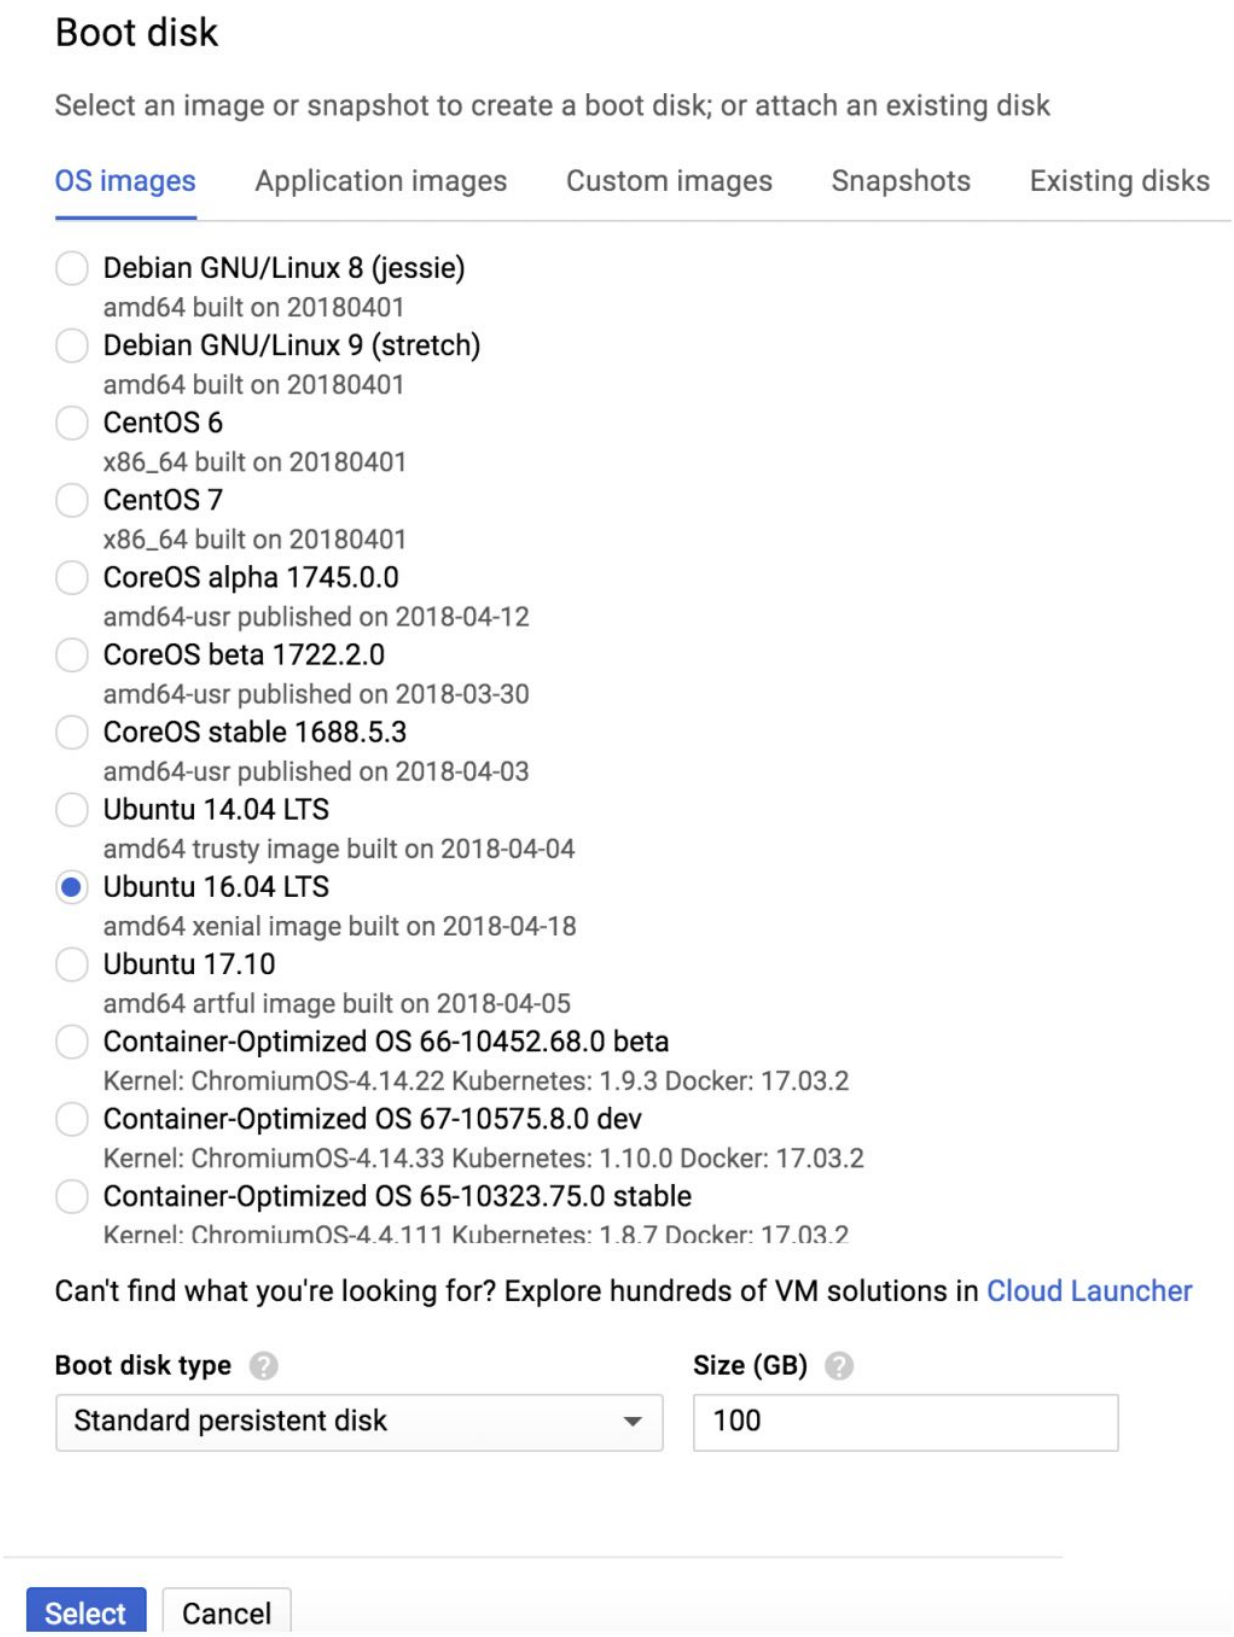

9. Select **Allow HTTP** and **Allow HTTPS Traffic**.

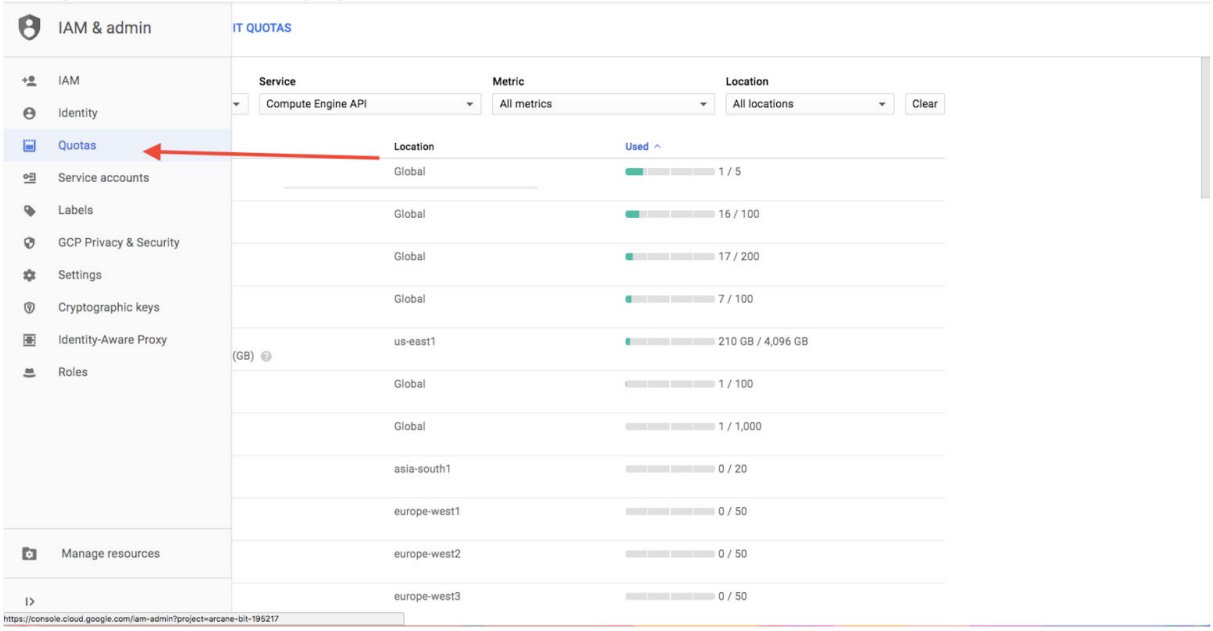

10. Create and launch your instance.

11. Run the following script, waiting for it to complete before proceeding.

curl -O http://developer.download.nvidia.com/compute/cuda/repos/ubuntu1604/x86_64/cuda-repo-ubuntu1604_8.0.61-1_amd64.deb

sudo dpkg -i ./cuda-repo-ubuntu1604_8.0.61-1_amd64.deb
sudo apt-get update
sudo apt-get -y install cuda-8-0

12. Run the following script:
sudo nvidia-smi -pm 1
echo 'export CUDA_HOME=/usr/local/cuda' >> ~/.bashrc
echo 'export PATH=$PATH:$CUDA_HOME/bin' >> ~/.bashrc
echo 'export LD_LIBRARY_PATH=$CUDA_HOME/lib64' >> ~/.bashrc
source ~/.bashrc

wget https://www.dropbox.com/s/44q4pacpv1otp7c/cudnn-8.0-linux-x64-v5.1.tgz
tar xzvf cudnn-8.0-linux-x64-v5.1.tgz
sudo cp cuda/lib64/* /usr/local/cuda/lib64/
sudo cp cuda/include/cudnn.h /usr/local/cuda/include/
sudo apt-get install python-dev python-pip libcupti-dev

curl https://bootstrap.pypa.io/get-pip.py

sudo pip install --upgrade https://storage.googleapis.com/tensorflow/linux/gpu/tensorflow_gpu-1.2.0-cp27-none-linux_x86_64.whl

sudo pip3 install tensorflow
sudo pip3 install keras
sudo jupyter notebook --ip 0.0.0.0 --port 8888 --allow-root

13. Proceed like in the other GCP tutorial.